


# MicroMQ

MCS Lab 2016

Marco Molteni  
Version 0.3, 2016-06-02

# 1. Introduction

Implement MicroMQ, a *minimal* pub/sub message broker in Erlang.



The clients can be simple [telnet](#) clients, while the server will be an Erlang program.

In any case, for testing, you will want to use an Erlang client :-)

## 2. Logistics

Small groups of 3 to 4 people.

## 3. Protocol

### 3.1. Connection

On connection, the broker will send to the client its ID as follows:

```
client_id: <client id>
```

(an unsigned integer)

### 3.2. Publish

To allow to use a [telnet](#) client, we will use a simple text-based protocol as follows:

```
topic: <topic name>  
body: <msg body>
```

where the message framing is expressed by two consecutive newlines (as in HTTP).

To which the broker will respond:

```
accepted: <topic name>
```

Topic creation is automatic: there is no need for the publisher to declare beforehand a topic: the first message sent on a topic (see example above) will also create the topic on the server side.

### 3.3. Subscribe

On the other hand, the subscriber will need to declare its interest explicitly:

```
subscribe: <topic1> [ , <topic2>, ... ]
```

to which the broker will respond:

```
subscribed: <topic1> [ , <topic2>, ... ]
```

If a client subscribes to a topic unknown to the broker, the broker will in any case accept the subscription and act upon it on the first message to the given topic.

### 3.4. Wildcard topic

There is a special topic, the wildcard: `*`. Subscribing to such topic will ask the broker to forward *all* messages received.

### 3.5. Forwarding

Upon reception of a message, the broker will forward it to all clients subscribed to the given topic (and to all the wildcard subscribers) as follows:

```
from: <client id>  
topic: <topic name>  
body: <msg body>
```

### 3.6. Status

When the client sends the command

```
status
```

the broker will respond with a summary of the status associated to the client:

```
status for client <client id>:  
subscribed topics: <list of subscribed topics>  
published topics: <list of published topics>  
received messages: <number of received messages>  
sent messages: <number of sent messages>
```

## 4. Configuration

The broker will have the following configuration parameters:

- TCP address (default: localhost only)
- TCP port (default: 5017)
- Maximum number of topics (default: 10'000)
- Maximum number of clients (default: 10'000)

## 5. Data structures

For the first time we are considering *big* data structures, and we also would like to survive a process crash.

Instead of keeping everything in the process state (probably a map or list in the [loop](#)), consider using an ETS table.

More information about ETS tables at:

- [LYSE: ETS](#)
- [ETS reference documentation](#)

## 6. Text manipulation

Remember that a socket should send and receive text as binary:

```
<<"topic: <topic name>\nbody: <msg body>\n\n">>
```

## 7. Documentation

Same short "rapport" as for PCAP and DRUM.

## 8. Build, test and run

The code must be buildable and eunit-testable with rebar, and must respect standard OTP directory layout: BEAM files go below [ebin/](#), not below [src/](#).

Compile without warnings:

```
$ rebar compile
```

Run tests without failures:

```
$ rebar eunit
```

Start Erlang (see also slide deck #2):

```
$ erl -pa ebin
```

The minimal rebar configuration *must* be (as the one in the MCS repo):

```
% Options for the Erlang compiler.
{erl_opts, [
    % DO NOT DISABLE THIS. If you get a warning, it means you are doing something wrong.
    warnings_as_errors
]}.

% Enable code coverage and generating the HTML details.
{cover_enabled, true}.

% Enable printing the coverage summary at the end of the build.
{cover_print_enabled, true}.
```

## 9. API

Please remember to respect the API.

Module containing the API: [micromq](#).

You can then add as many implementation modules as you see fit.

API

Start the broker with default values:

```
-spec start_link() -> ok | {error, Reason}.
```

Start the broker with custom values:

```
-type option() ::
    {address, Address} |
    {port, Port} |
    {max_topics, N} |
    {max_clients, N}.

-spec start_link(Options) -> ok | {error, Reason} when
    Options :: [option()].
```

Stop the broker. Since it has arity 0, you will need a registered process:

```
-spec stop() -> ok.
```

## 10. Spawning and socket ownership

Depending how you choose to spawn a process to server the socket, you might need to call `gen_tcp:controlling_process(Socket, Pid)`. Please refer to <http://learnyousomeerlang.com/buckets-of-sockets> and search for the text "controlling\_process" for details.

## 11. Optional

- Implement the [system](#) topic (see below).
- Implement [zeroconf](#) autodiscovery of the broker.
- Implement message persistence (see below).

### 11.1. System topic

The broker offers a special topic, `system`. If a client sends a message to topic `system`, the broker will respond:

```
error: system is a reserved topic
```

Any client can subscribe to the system topic.

The broker will send, each 20 seconds (configurable), a message on the system topic as follows

```
from: broker
topic: system
body:
  system uptime: <system uptime in seconds>
  connected clients: <number of connected clients>
  subscribed topics: <list of subscribed topics>
  published topics: <list of published topics>
  received messages: <number of received messages>
  sent messages: <number of sent messages>
```

The period of the `system` topic is configurable by passing the following option to [start\\_link/2](#):

```
{system_topic_period, N}
```

### 11.2. Message persistence

If a client, before disconnecting, sends the command

```
persist: <key>
```

to which the broker will respond

```
persisted: <key>
```

the server will associate `key` to the current `client id` and store all messages that match the client subscriptions.

When the client reconnects and sends the command

```
restore: <key>
```

the broker will replay to it all the stored messages.

Note: consider how to optimize the case when multiple clients send the `persist` command and are subscribed to the same topics: what do you do? Do you copy N times the same message in N persistence queues or can you do something better?

If message persistence is implemented, the `system` topic will have an additional information:

```
persisted clients: <N>  
persisted messages: <M>
```